Held-Out Technical Note: Scaled Matrix Products

When matching continuous queries against a memory bank of keys, the inner product magnitude scales proportionally to the dimensionality of the vector space.

To preserve gradient dynamics across deep layers, scores are divided by the square root of the projection dimension before softmax normalization.

In autoregressive generation, future token positions must be strictly masked to prevent information leakage across time steps.